



Data Handling: Import, Cleaning and Visualisation

Lecture 10:

Data Analysis and Basic Statistics with R

Prof. Dr. Ulrich Matter

05/12/2019

Recap: Data Preparation

The dataset is imported, now what?


- In practice: still a long way to go.
- Parsable, but messy data: Inconsistencies, data types, missing observations, wide format.
- **Goal** of data preparation: Dataset is ready for analysis.
- **Key conditions:**
 1. Data values are consistent/clean within each variable.
 2. Variables are of proper data types.
 3. Dataset is in 'tidy' (in long format)!

Some vocabulary

Following Wickham (2014):

- **Dataset**: Collection of **values** (numbers and strings).
- Every value belongs to a **variable** and an **observation**
- **Variable**: Contains all values that measure the same underlying attribute across units.
- **Observation**: Contains all values measured on the same unit (e.g., a person).

Tidy data



Tidy data. Source: Wickham and Grolemund (2017), licensed under the [Creative Commons Attribution-Share Alike 3.0 United States](https://creativecommons.org/licenses/by-sa/3.0/) license)

Data Analysis with R

Merging (Joining) datasets

- Combine data of two datasets in one dataset.
 - Why?
- Needed: Unique identifiers for observations ('keys').

Merging (joining) datasets: example

```
# load packages
```

```
library(tidyverse)
```

```
# initiate data frame on persons personal spending
```

```
df_c <- data.frame(id = c(1:3,1:3),  
                  money_spent= c(1000, 2000, 6000, 1500, 3000, 5500),  
                  currency = c("CHF", "CHF", "USD", "EUR", "CHF", "USD"),  
                  year=c(2017,2017,2017,2018,2018,2018))
```

```
df_c
```

```
##   id money_spent currency year  
## 1  1      1000      CHF 2017  
## 2  2      2000      CHF 2017  
## 3  3      6000      USD 2017  
## 4  1      1500      EUR 2018  
## 5  2      3000      CHF 2018  
## 6  3      5500      USD 2018
```


Merging (joining) datasets: example

```
# initiate data frame on persons' characteristics
df_p <- data.frame(id = 1:4,
                   first_name = c("Anna", "Betty", "Claire", "Diane"),
                   profession = c("Economist", "Data Scientist", "Data Scientist", "Economist"),
                   stringsAsFactors = FALSE)

df_p
```

```
##   id first_name    profession
## 1  1      Anna    Economist
## 2  2     Betty Data Scientist
## 3  3    Claire Data Scientist
## 4  4     Diane    Economist
```

Merging (joining) Datasets: Example

```
df_merged <- merge(df_p, df_c, by="id")  
df_merged
```


```
##   id first_name      profession money_spent currency year  
## 1  1      Anna      Economist      1000      CHF 2017  
## 2  1      Anna      Economist      1500      EUR 2018  
## 3  2      Betty Data Scientist      2000      CHF 2017  
## 4  2      Betty Data Scientist      3000      CHF 2018  
## 5  3      Claire Data Scientist      6000      USD 2017  
## 6  3      Claire Data Scientist      5500      USD 2018
```

Merging (joining) datasets: concept

x		y	
1	x1	1	y1
2	x2	2	y2
3	x3	4	y3

Join setup. Source: Wickham and Grolemund (2017), licensed under the [Creative Commons Attribution-Share Alike 3.0 United States](#) license.


Merging (joining) datasets: concept



Join setup. Source: Wickham and Grolemund (2017), licensed under the [Creative Commons Attribution-Share Alike 3.0 United States](https://creativecommons.org/licenses/by-sa/3.0/) license.

Merging (joining) datasets: concept


Merge: Inner join



Inner join. Source: Wickham and Grolemund (2017), licensed under the [Creative Commons Attribution-Share Alike 3.0 United States](#) license.

Merging (joining) datasets: concept


Merge all x: Left join



Outer join. Source: Wickham and Grolemund (2017), licensed under the [Creative Commons Attribution-Share Alike 3.0 United States](#) license.

Merging (joining) datasets: concept


Merge all y: Right join



Outer join. Source: Wickham and Grolemund (2017), licensed under the [Creative Commons Attribution-Share Alike 3.0 United States](https://creativecommons.org/licenses/by-sa/4.0/) license.


Merging (joining) datasets: concept

Merge all x and all y: Full join



Outer join. Source: Wickham and Grolemund (2017), licensed under the [Creative Commons Attribution-Share Alike 3.0 United States](#) license.

Merging (joining) datasets: concept



Join Venn Diagramm. Source: Wickham and Grolemund (2017), licensed under the [Creative Commons Attribution-Share Alike 3.0 United States](https://creativecommons.org/licenses/by-sa/3.0/) license.

Merging (joining) datasets: R

Overview by Wickham and Grolemund (2017):

dplyr (tidyverse)

base::merge

inner_join(x, y)

merge(x, y)

left_join(x, y)

merge(x, y, all.x = TRUE)

right_join(x, y)

merge(x, y, all.y = TRUE),

full_join(x, y)

merge(x, y, all.x = TRUE, all.y = TRUE)

Data Summaries: Selecting, Filtering, and Mutating

Data summaries

- First step of analysis.
- Get overview over dataset.
- Show key aspects of data.
 - Inform your own statistical analysis.
 - Inform audience (helps understand advanced analytics parts)

Data summaries: first steps

- **Select** subset of variables (e.g., for comparisons).
- **Filter** the dataset (some observations not needed in **this** analysis).
- **Mutate** the dataset: additional values needed

Select, filter, mutate in R (tidyverse)

- `select()`
- `filter()`
- `mutate()`

Data Summaries: Aggregate Statistics

Descriptive/aggregate statistics

- Overview of key characteristics of main variables used in analysis.
- Key characteristics:
 - mean
 - standard deviation
 - No. of observations
 - etc.

Aggregate statistics in R

1. Function to compute statistic (`mean()`).
2. Function to **apply** the statistics function to one or several columns in a tidy dataset.
 - All values.
 - By group (observation categories, e.g. by gender)

Aggregate statistics in R

- `summarise()` (in `tidyverse`)
- `group_by()` (in `tidyverse`)
- `sapply()`, `apply()`, `lapply()`, etc. (in `base`)

Understanding Statistics and Probability with Code

Random numbers and computation

Can computers generate random numbers?!

Random draws and distributions

```
normal_distr <- rnorm(1000)  
hist(normal_distr)
```

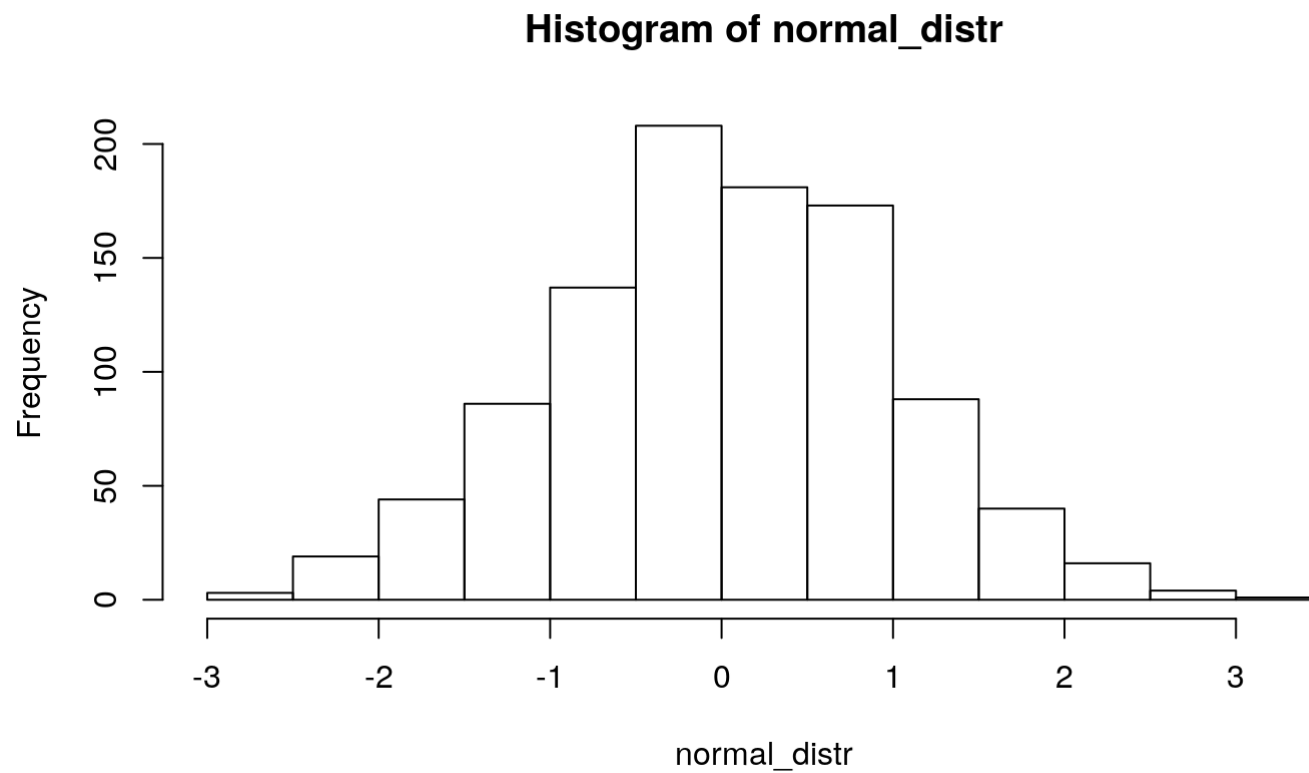


Illustration of variability

```
# draw a random sample from a normal distribution with a large standard deviation  
largevar <- rnorm(10000, mean = 5000, sd = 5)  
# draw a random sample from a normal distribution with a small standard deviation  
littlevar <- rnorm(10000, mean = 5000, sd = 1)
```

Illustration of variability

```
# visualize the distributions of both samples with a density plot  
plot(density(littlevar), col = "blue",  
      xlim=c(min(largevar), max(largevar)), main="Income Distribution")  
lines(density(largevar), col = "red")
```

Skewness and kurtosis

```
# install the R-package called "moments" with the following command (if not installed)  
# install.packages("moments")  
  
# load the package  
library(moments)
```


Skewness

- Skewness refers to how symmetric the frequency distribution of a variable is.
- For example, a distribution can be **'positively skewed'** meaning it has a **long tail on the right** and a lot of 'mass' (observations) on the left.

Skewness: R example

```
# draw a random sample of simulated data from a normal distribution  
# the sample is of size 1000 (hence, n = 1000)  
sample <- rnorm(n = 1000)  
  
# plot a histogram and a density plot of that sample  
# note that the distribution is neither strongly positively nor negatively skewed  
# (this is to be expected, as we have drawn a sample from a normal distribution!)  
hist(sample)
```

Skewness: R example

```
# now compute the skewness  
skewness(sample)
```

```
## [1] -0.06196048
```


Skewness: R example

```
# now we intentionally change our sample to be strongly positively skewed  
# we do that by adding some outliers (observations with very high values) to the sample  
sample <- c(sample, (rnorm(200) + 2), (rnorm(200) + 3))  
  
# have a look at the distribution and re-calculate the skewness  
hist(sample)
```

Skewness: R example

have a look at the distribution and re-calculate the skewness

```
hist(sample)
```



Kurtosis

- Kurtosis refers to how much 'mass' a distribution has in its 'tails'.
- Tells us something about whether a distribution tends to have a lot of outliers.

Kurtosis: R example

```
# draw a random sample of simulated data from a normal distribution  
# the sample is of size 1000 (hence, n = 1000)  
sample <- rnorm(n = 1000)  
  
# plot the density & compute the kurtosis  
hist(sample)
```

Kurtosis: R example

```
# compute the kurtosis  
kurtosis(sample)
```


```
## [1] 3.144589
```


Kurtosis: R example

```
# now lets remove observations from the extremes in this distribution  
# we thus intentionally alter the distribution to have less mass in its tails  
sample <- sample[ sample > -0.6 & sample < 0.6]
```

Kurtosis: R example

plot the distribution again and see how the tails have changed
`hist(sample)`



Compute the skewness in R

Skewness

own implementation

```
sum((sample-mean(sample))^3) / ((length(sample)-1) * sd(sample)^3)
```

```
## [1] -0.09341458
```

implementation in moments package

```
skewness(sample)
```

```
## [1] -0.09351786
```

Compute the kurtosis in R

Kurtosis

own implementation

```
sum((sample-mean(sample))^4) / ((length(sample)-1) * sd(sample)^4)
```

```
## [1] 1.806926
```

implementation in moments package

```
kurtosis(sample)
```

```
## [1] 1.810923
```

The Law of Large Numbers (LLN)

- Important statistical property.
- Essentially describes how the behavior of **sample averages** is related to **sample size**.
- States that the **sample mean** can come arbitrarily close to the **population mean** by increasing the sample size N .

The Law of Large Numbers (LLN): playing dice

- Roll a fair die, record result: either 1, 2, 3, 4, 5, or 6.
- Probability of each possible outcome is $1/6$.

The Law of Large Numbers (LLN): playing dice

- Roll a fair die, record result: either 1, 2, 3, 4, 5, or 6.
- Probability of each possible outcome is $1/6$.
- Expected value (average in the long run): $(1 + 2 + 3 + 4 + 5 + 6)/6 = 3.5$

The Law of Large Numbers (LLN): playing dice

- Roll a fair die, record result: either 1, 2, 3, 4, 5, or 6.
- Probability of each possible outcome is $1/6$.
- Expected value (average in the long run): $(1 + 2 + 3 + 4 + 5 + 6)/6 = 3.5$
- Proof?
 - Mathematically.
 - Or: Experiment/Simulation (with R).

LLN in R

```
# first we define the potential values a die can take  
dvalues <- 1:6 # the : operator generates a regular sequence of numbers (from:to)  
dvalues
```

```
## [1] 1 2 3 4 5 6
```

```
# define the size of the sample n (how often do we roll the die...)  
# for a start, we only roll the die ten times  
n <- 10  
# draw the random sample: 'roll the die n times and record each result'  
results <- sample( x = dvalues, size = n, replace = TRUE)  
# compute the mean  
mean(results)
```

```
## [1] 3.9
```

LLN in R

```
n <- 100
# draw the random sample: 'roll the die n times and record each result'
results <- sample( x = dvalues, size = n, replace = TRUE)
# compute the mean
mean(results)

## [1] 3.48
```

LLN in R

```
# essentially, what we are doing here is repeating the experiment above many times,  
# each time increasing n  
# define the set of sample sizes  
ns <- seq(from = 10, to = 10000, by = 10)  
# initiate an empty list to record the results  
means <- list()  
length(means) <- length(ns)  
# iterate through each sample size: 'repeat the die experiment for each sample size  
for (i in 1:length(ns)) {  
  
    means[[i]] <- mean(sample( x = dvalues, size = ns[i], replace = TRUE))  
  
}
```

LLN in R

```
# visualize the result: plot sample means against sample size
plot(ns, unlist(means),
     ylab = "Sample Mean",
     xlab = "Sample Size",
     pch = 16,
     cex = .6)
abline(h = 3.5, col = "red")
```

Q&A

References

Wickham, Hadley. 2014. "Tidy Data." **Journal of Statistical Software, Articles** 59 (10): 1–23.
<https://doi.org/10.18637/jss.v059.i10>.

Wickham, Hadley, and Garrett Golemund. 2017. Sebastopol, CA: O'Reilly. <http://r4ds.had.co.nz/>.